

Parallel CRC Algorithm and Implementation with CUDA

Mirco De Marchi - VR445319 Luigi Capogrosso - VR445456

Abstract—*Cyclic redundancy checks (CRCs) form a powerful class of codes suited especially for the detection of burst errors in data storage and communication applications. The use of cyclic redundancy codes, referred to more briefly in this document as check words, is one of the most well known error checking mechanism. It is customarily placed at the end of a data block, occupying one or more bytes of the transmitted message. In this paper, we present a Parallel Cyclic Redundancy Check (PCRC) algorithm improved with CUDA. Our approach is to systematically decompose the original input message into a set of subsequences based on the Galois field theory.*

I. STATE OF THE ART

Devices such as computers, routers, networking equipment, and components that communicate information internally and/or with other devices. For example, computers might communicate across a local area network (LAN) using Ethernet protocol, or application-specific integrated circuits (ASICs) may communicate with each other over a single or parallel bit bus. **It is important for these devices to reliably communicate and detect errors in their communication.**

One common technique for detecting transmission errors is a technique known as **cyclic redundancy check (CRC)**. A CRC allows errors detection using only a small number of redundant bits typically sent along with the communicated information. For example, a 32-bit CRC gives strong protection against burst bits errors in messages that are thousands bits long.

The fundamental principle upon which CRC is based can be expressed equivalently in one of three ways.

- 1) CRC can be described in terms of binary numbers division;
- 2) It can be performed utilizing a polynomials division;
- 3) The utility in implementing CRCs is often realized by designing special check circuits in which Exclusive Or (XOR) and other elementary binary operands generate partial results that can be used in CRC algorithm;

CRCs are specifically designed to protect against common types of errors on communication channels, where they can provide quick and reasonable assurance of the integrity of the messages delivered. However, CRCs are not suitable for protecting against intentional alteration of data.

II. INTRODUCTION

Modern day technology permits a large amounts of data transfer between computer or communication systems. During transmission these bits may become altered, thereby creating an error. In order for one computer system to detect when

data sent by another computer system has an error, certain status information is often transmitted along with the data. This status information may include parity (even or odd), longitudinal redundancy check character, or cyclic redundancy check code.

CRCs are based on the theory of **systematic cyclic codes**, which encode messages by adding a fixed-length check value. Cyclic codes are not only simple to implement but **have the benefit of being particularly well suited for the detection of burst errors**: contiguous sequences of erroneous data symbols in messages.

Specification of a CRC code requires definition so-called **generator polynomial**. For the purposes of constructing polynomial codes, we identify a string of n symbols a_{n-1}, \dots, a_0 with the polynomial:

$$a_{n-1}x^{n-1} + \dots + a_1x + a_0$$

This polynomial becomes the **divisor** in a **polynomial long division**, which takes the message as the **divident**, the **quotient** is discarded and the **remainder** becomes the result. The important is that **the polynomial coefficients are calculated according to the arithmetic of a finite field**, so the addition operation can always be performed bitwise-parallel (there is no carry between digits). A finite field is a set on which the operations of multiplication, addition, subtraction and division are defined and satisfy certain basic rules.

A CRC is called an N -bit CRC when its check value is N bits long. The simplest error-detection system, the parity bit, is in fact a 1-bit CRC: it uses the generator polynomial $x + 1$ (two terms), and has the name CRC-1.

A CRC-enabled device calculates a short, fixed-length binary sequence, known as CRC, that will be appended to the end of message sent. The receiving device recalculates the CRC code over the received message payload and If the CRC values do not match, then the message contains a data error. The device may take corrective action, such as rereading the block or requesting that it be sent again. Otherwise, the data is assumed to be error-free.

The common hardware solution is the “Linear Feedback Shift Register” (LFSR), which is a simple bitwise architecture for both encoding and decoding the message. This approach typically calculates the CRC for a N -bit message in N clock cycles. This approach is not efficient at high bit rates.

III. BACKGORUND

The CRC see the message A of length n as a polynomial $A(x)$ of degree $n - 1$ in which every bit of message is the

coefficient of the respective monomial:

$$A = [a_{n-1}, a_{n-2}, \dots, a_2, a_1, a_0]$$

$$A(x) = a_{n-1}x^{n-1} + a_{n-2}x^{n-2} + \dots + a_2x^2 + a_1x + a_0$$

For example, the input data 0x25=0010 0101 is taken as:

$$A(x) = 0x^7 + 0x^6 + 1x^5 + 0x^4 + 0x^3 + 1x^2 + 0x^1 + 1x^0$$

Then for the CRC computation is used a polynomial generator $G(x)$ that is polynomial of degree m . The major index coefficient is always 1, in this way the polynomial generator guarantees to always be of m degree.

In a reverse way than the message A , the polynomial generator can be represented as a sequence of bit G of size $m + 1$ with the most significant bit always to 1. This the reason why you will always find the value representation of the polynomial generator G of size m , because the bit 1 to the left is implicit.

$$G = [g_m, g_{m-1}, \dots, g_1, g_0]$$

$$G = g_mx^m + g_{m-1}x^{m-1} + \dots + g_1x + g_0$$

For example, Ethernet uses the following 32-bit polynomial value:

$$G(x) = 1 + x + x^2 + x^4 + x^5 + x^7 + x^8 + x^{10} + x^{11} + x^{12} + x^{16} + x^{22} + x^{23} + x^{26} + x^{32}$$

The **value of the polynomial generator** has a **great impact on its error detection** capabilities, its design is non-trivial and requires serious math knowledge.

The CRC detection code is the result of the remainder after dividing the original message A concatenated with m zero bits by a polynomial generator in binary modulo 2 arithmetic.

In the polynomial version this definition of CRC is equivalent to:

$$CRC[A(x)] = A(x)x^m \bmod 2(G(x))$$

The binary modulo 2 arithmetic makes every operation between bits independent between each other. This means that **the carry bit can be completely forgotten**. In practice the arithmetic operations are identical but the carryover is ignored. In particular the binary modulo 2 remainder of the division always generates a result that is a degree less of the degree of the divisor. Therefore the CRC result will always be a polynomial of degree $m - 1$ that can be viewed as a binary code of size m . In fact:

$$CRC[A(x)] = A(x)x^m \bmod 2(G(x))$$

is equal to:

$$\text{degree}(CRC[A(x)]) = m - 1$$

This because the main difference between different CRC code is the length m of the code generated. The most important CRC Implementations are **CRC8**, **CRC16**, **CRC32** and **CRC64**. The longer the length of the CRC code, the more likely bit errors shall be detected. In general if a network protocol provides very long message should be used a longer

CRC as CRC32 or CRC64, if a network exchanges short message can be used shorter CRC error detection code as CRC8 or CRC16.

Then there are different parametrization of the CRC:

- **Polynomial generator:** normal, reversed (from the least significant bit to the most), reciprocal (reversed but respect the polynomial coefficients), reversed and reciprocal in combination and the Koopman representation (discarding the least significant bit);
- **Initial value of the CRC:** in most algorithms is 0, in some cases is set with all bit to 1, but in general could be any value;
- **Input reflected:** each input byte can be used in reverse order before being used in the calculation;
- **Output reflected:** the result can be returned in reverse order over the whole CRC value;
- **Final XOR value:** is value xored to the final CRC value before being returned and, in case, after the output reflected step.

IV. LEMMAS

Before we present our parallel CRC algorithm, we need the following lemma, which gives the well-known congruence properties for polynomial division.

$$(A(x) + B(x)) \bmod G(x) = A(x) \bmod G(x) + B(x) \bmod G(x) \quad (1)$$

$$A(x)B(x) \bmod G(x) = ((A(x) \bmod G(x))(B(x) \bmod G(x))) \bmod G(x) \quad (2)$$

There is a system of arithmetic known as **Galois field arithmetic**. A Galois field is an example of algebraic field, which has a set of elements called **numbers or field elements**, and a definition of two operations called **addition and multiplication** such that the formal properties satisfied by the real arithmetic system (e.g. commutativity, associativity, and distributivity) are satisfied.

For each positive integer M , there is a Galois field called $GF(2^M)$ that has 2^M elements in it. The set of elements of $GF(2^M)$ can be represented as the set of M -bit binary numbers in the range from 0 to $2^M - 1$. The addition over Galois Field $GF(2^M)$ is defined as bit-by-bit modulo-2 addition which is equivalent to exclusive-or (i.e., XOR) operation.

In practice, all commonly used CRCs employ the Galois field of two elements, $GF(2)$. The two elements are usually called 0 and 1, comfortably matching computer architecture. Multiplication of two elements a and b in $GF(2^M)$ will produce an element $c = ab$ by polynomial multiplication modulo the irreducible polynomial $G(x)$:

$$C(x) = A(x)B(x) \bmod G(x)$$

where $A(x)$, $B(x)$ and $C(x)$ is the polynomial form of a , b , and c , respectively.

V. CRC ALGORITHMS

Common CRC algorithms are bitwise and byte-wise.

In the bitwise CRC algorithm, 1 input bit is processed at a time as the name indicates, using long division.

First we append M zeros (M is the number of bits in CRC) to the original N -bit message to the least significant bit (LSB) and define the appended message as a running message.

Second, check the most significant bit (MSB) of the running message.

If the MSB of the running message is 1, subtract the generator polynomial from the M most significant bits of this running message and shift the result to left by 1 bit and store the result to the running message. Otherwise, (i.e., the MSB of the running message is 0), just shift the running message to left by 1 bit and store the result to the running message.

The second step is repeated N times and at the end the remainder will be the M most significant bits, which is the CRC. Note that in bitwise CRC algorithm, we need to perform N times of operation where in each operation we need to perform checking the MSB bit, modulo-2 subtraction (conditionally), and left shift.

This is a **pseudocode of the bitwise CRC algorithm implementation:**

```
// Result size: N + M.
append(message, M)

// Usually 0.
message[M-1 : 0] = crc_initial_value

// Check each bit of the message.
for i in range(0 : N)
{
    LSB = (N + M - 1) - i
    if (message[LSB] == 1)
    {
        message[LSB : LSB-(M-1)] ^=
            polynomial_generator;
    }
}

return message[M-1 : 0];
```

In byte-wise CRC algorithm, one byte is checked at a time.

In the first step, similar to bitwise CRC algorithm, we append M zeros to the original message after the LSB.

If the original message size N is not a multiple of 8, we need to pre-append a number of 0's (in the range of 1 to 7) before the MSB to make the appended message having size of multiple of 8. We define this appended message as a running message.

Second, perform table lookup based on the MSB 8 bits to find the M -bit remainder. This M -bit remainder will be XORed with the following MSB byte in the running message. Then we left shift the running message by 8 bits.

Repeat the second step by $\lceil N/8 \rceil$ times. In total there are $\lceil N/8 \rceil$ operations with $28 \times M$ bit of memory for table lookup. Note that the table has 256 entries each of which has M bits.

```
const uint8_t crc8_lu[] = {
    0x0,  0x1d, 0x3a, 0x27, 0x74, 0x69, 0x4e, 0x53,
    0xe8, 0xf5, 0xd2, 0xcf, 0x9c, 0x81, 0xa6, 0xbb,
    0xcd, 0xd0, 0xf7, 0xea, 0xb9, 0xa4, 0x83, 0x9e,
    0x25, 0x38, 0x1f, 0x2,  0x51, 0x4c, 0x6b, 0x76,
    0x87, 0x9a, 0xbd, 0xa0, 0xf3, 0xee, 0xc9, 0xd4,
    0x6f, 0x72, 0x55, 0x48, 0x1b, 0x6,  0x21, 0x3c,
    0x4a, 0x57, 0x70, 0x6d, 0x3e, 0x23, 0x4,  0x19,
    0xa2, 0xbf, 0x98, 0x85, 0xd6, 0xcb, 0xec, 0xf1,
    0x13, 0xe,  0x29, 0x34, 0x67, 0x7a, 0x5d, 0x40,
    0xfb, 0xe6, 0xc1, 0xdc, 0x8f, 0x92, 0xb5, 0xa8,
    0xde, 0xc3, 0xe4, 0xf9, 0xaa, 0xb7, 0x90, 0x8d,
    0x36, 0x2b, 0xc,  0x11, 0x42, 0x5f, 0x78, 0x65,
    0x94, 0x89, 0xae, 0xb3, 0xe0, 0xfd, 0xda, 0xc7,
    0x7c, 0x61, 0x46, 0x5b, 0x8,  0x15, 0x32, 0x2f,
    0x59, 0x44, 0x63, 0x7e, 0x2d, 0x30, 0x17, 0xa,
    0xb1, 0xac, 0x8b, 0x96, 0xc5, 0xd8, 0xff, 0xe2,
    0x26, 0x3b, 0x1c, 0x1,  0x52, 0x4f, 0x68, 0x75,
    0xce, 0xd3, 0xf4, 0xe9, 0xba, 0xa7, 0x80, 0x9d,
    0xeb, 0xf6, 0xd1, 0xcc, 0x9f, 0x82, 0xa5, 0xb8,
    0x3,  0x1e, 0x39, 0x24, 0x77, 0x6a, 0x4d, 0x50,
    0xa1, 0xbc, 0x9b, 0x86, 0xd5, 0xc8, 0xef, 0xf2,
    0x49, 0x54, 0x73, 0x6e, 0x3d, 0x20, 0x7,  0x1a,
    0x6c, 0x71, 0x56, 0x4b, 0x18, 0x5,  0x22, 0x3f,
    0x84, 0x99, 0xbe, 0xa3, 0xf0, 0xed, 0xca, 0xd7,
    0x35, 0x28, 0xf,  0x12, 0x41, 0x5c, 0x7b, 0x66,
    0xdd, 0xc0, 0xe7, 0xfa, 0xa9, 0xb4, 0x93, 0x8e,
    0xf8, 0xe5, 0xc2, 0xdf, 0x8c, 0x91, 0xb6, 0xab,
    0x10, 0xd,  0x2a, 0x37, 0x64, 0x79, 0x5e, 0x43,
    0xb2, 0xaf, 0x88, 0x95, 0xc6, 0xdb, 0xfc, 0xe1,
    0x5a, 0x47, 0x60, 0x7d, 0x2e, 0x33, 0x14, 0x9,
    0x7f, 0x62, 0x45, 0x58, 0xb,  0x16, 0x31, 0x2c,
    0x97, 0x8a, 0xad, 0xb0, 0xe3, 0xfe, 0xd9, 0xc4
};
```

Figure 1. LOOKUP TABLE USED IN CRC8 BYTEWISE

The table entries can be pre-computed since they only depend on the generator polynomial.

This is a **pseudocode of the byte-wise CRC algorithm implementation:**

```
// Result size: N + M.
append(message, M)

// Usually 0
message[M-1 : 0] = crc_initial_value

// Check each byte of the message.
for i in range(0:N / 8-1)
{
    LSB = (N + M - 1) - (i * 8);
    message[LSB-8 : LSB-8-(M-1)] ^=
        lookup_table[message[LSB : LSB-7]];
}

return message[M-1 : 0];
```

VI. THE CRC LOOKUP TABLE OPTIMIZATION

So far **the algorithm is quite inefficient as it works bit by bit**. For larger input data, this could be quite slow. But how can our CRC algorithm be accelerated?

The dividend is the current CRC byte value and a byte can only take 256 different values. The polynomial (divisor) is fixed. **Why not precompute the division for each possible byte by the fixed polynomial and store these result in a lookup table?** This is possible as the remainder is always the same for the same dividend and divisor! Then the input stream can be processed byte by byte instead of bit by bit.

In computer science, a **lookup table** is an array that **replaces runtime computation with a simpler array indexing operation** (Figure 1). The savings in terms of processing time can be significant, since retrieving a value from memory is often faster than undergoing an “expensive” computation or input/output operation. The tables may be precalculated and stored in static program storage, calculated (or “pre-fetched”) as part of a program’s initialization phase (memoization), or even stored in hardware in application-specific platforms.

To do a table lookup (index an array), the result needs to be right shifted, or the high order byte of the CRC can be right shifted first, then XOR’ed with a byte of data to produce the new intermediate high order byte of the CRC, then that has to be cycled 8 times to produce the result of cycling the CRC 8 time, or this can be done in advance for all 256 possible 8 bit values, to generate a table, and a table lookup can be used instead of cycling 8 times. After this cycling, the low order byte will need to be shifted left, so you end up with the code shown.

A left shifting CRC is similar to doing a long hand division in binary, using the CRC polynomial as the divisor, and the data as a long dividend, except that XOR is used instead of subtraction, and the final remainder is the CRC. It turns out that XOR’ing 8 data (dividend) bits at a time doesn’t affect the result, since each quotient bit is based only on the most significant bit of CRC polynomial and the current most significant bit of the working dividend (data), which allows for the optimization of using a table lookup.

A simple pseudocode to generate the lookup table could be:

```
for i in range(0 : 2{8} - 1)
{
    lookup_table[i] = crc_bitwise(i)
}
```

VII. PARALLEL CRC ALGORITHM

In this section we present our “**Parallel Cyclic Redundancy Check**” (PCRC) algorithm that performs CRC computation for any length of message in parallel.

For a given message with any length, we first chunk the message into blocks, each of which has a fixed size equal to the degree of the generator polynomial. Then we perform CRC computation among the chunked blocks in parallel using Galois Field Multiplication and Accumulation (GFMAC).

The message is seen as divided in M bit size chunk and the following is its polynomial representation:

$$A(x) = W_{n-1}(x)x^{(n-1)M} + \dots + W_1(x)x^M + W_0(x)$$

Where each $W_i(x)$ polynomial is a chunk of the message. From this equation and the CRC definition, we can compute the CRC for the chunked message by:

$$CRC[A(x)] = W_{n-1}x^{nM} \bmod G(x) + \dots + W_0x^M \bmod G(x)$$

Furthermore, from the Galois Field operations lemma, we obtain that:

$$W_i(x)x^{(i+1)M} \bmod G(x) = (W_i(x) \bmod G(x))x^{(i+1)M} \bmod G(x)$$

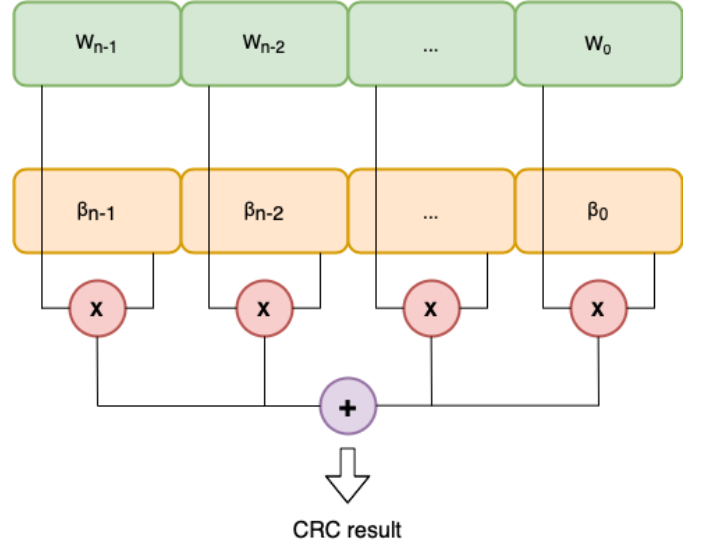


Figure 2. ILLUSTRATION OF PARALLEL CRC ALGORITHM OVER CHUNK OF M BITS

The degree of the polynomial $W_i(x)$ for each chunk is $M-1$, therefore it is less than M , it means that $W_i(x) \bmod G(x) = W_i(x)$. On the other hand the beta coefficients are defined as $\beta_i = x^{(i+1)M} \bmod G(x)$ for $i = 0, 1, \dots, n-1$. Therefore we have:

$$CRC[A(x)] = W_{n-1} \otimes \beta_{n-1} \oplus \dots \oplus W_0 \otimes \beta_0$$

Note that the operations \otimes, \oplus in the above equation is Galois Field multiplication, addition over $GF(2^M)$, respectively. The following algorithm is illustrated in Figure 2.

- 1) Load N -bit message in the chunked form of (W_{n-1}, W_{n-2}, W_0) where each chunk is M bits. (Note: $N = nM$).
- 2) Initially setup the generator polynomial $G(x)$ and its degree M . In the same time, pre-compute the beta factors $(\beta_{n-1}, \beta_{n-2}, \beta_0)$ which depend only on the polynomial and its degree.
- 3) Perform n -pair Galois field multiplications in parallel and then XOR the products. **This generates the CRC results.**

To compute the beta factors, we have:

$$\beta_0 = x^M \bmod G(x) = [g_{m-1}, \dots, g_0]$$

$$\beta_1 = x^{2M} \bmod G(x) = \beta_0 \otimes \beta_0 = \beta_0^2$$

...

$$\beta_{n-1} = x^{nM} \bmod G(x) = \beta_0^n$$

Note that beta factors $(\beta_{n-1}, \beta_{n-2}, \beta_0)$ do not depend on the incoming message but only on the generator polynomial. Since the polynomial is static, this should be pre-computed once and used repeatedly in the CRC calculation.

A simple pseudocode to generate the beta array could be:

```
for i in range(0 : N)
{
    shift_buffer = 1 << (M * (i + 1));
```

```

    beta[i]] = mod2(shift_buffer,
                    generator_poly);
}

```

Each thread is executed over a M bit size chunk of the original message and performs the following operations:

```

// Get the data of this thread.
W_i = original_message[global_index];
beta_i = beta[global_index];

// Perform binary modulo 2 multiplication.
mul = mod2_mul(W_i, beta_i);

// Perform binary modulo 2 reminder.
mod = mod2_mod(mul, generator_poly);

// Copy in shared memory.
shared_memory[threadX] = mod;
sync();

// XOR all data in shared memory.
return xored(shared_memory);

```

The last operation, that performs the Galois Field addition, xoring all the results of the Galois Field multiplication calculated from each thread, has been done in 2 ways:

- 1) The standard implementation takes the first thread, stop all the other and performs the XOR operation on all shared memory. This solution has a problem of divergence because one thread is in execution while the other has nothing to do.
- 2) The reduction implementation takes the first half of all thread and performs iteratively the XOR operation between a pair of values in shared memory. In this way the workload is more spread across all threads and there is less divergence.

Finally, since the thread execution is done in blocks, the result of the device execution will be an array of partial results. Each value of the array of partial results will be the xored value of all the Galois Field multiplication of the message chunk of that specific thread block. Therefore the host has to perform the XOR operation over the array of partial results to find the right CRC code. The Figure ?? shows how the CRC result is obtained by the xoring computation over the thread blocks.

VIII. PERFORMANCE ANALYSIS

Theoretically this level of performance is hundreds of times faster than bitwise CRC algorithm or tens of times faster than bitwise parallel CRC algorithm.

With our software implementation we have achieved the following results:




Figure 3. XOR OPERATIONS PERFORM BY DEVICE IN RELATION WITH THE ONE PERFORM BY HOST

A. $N = 2^{16}$. $BlockSize = 64$. $StreamDim = 4$.
 $SegSize = N/StreamDim$

Test PCRC 64 blocksize	Speedup
PCRC8	34.71340
PCRC8 with reduction	37.40804
PCRC8 with task parallelism	27.86828
PCRC8 bitwise comparison	2.58712
PCRC8 bitwise comparison with reduction:	2.61057
PCRC8 bitwise comparison with task parallelism	1.93434
PCRC16	34.98040
PCRC16 with reduction	40.59672
PCRC16 with task parallelism	26.11727
PCRC16 bitwise comparison	3.44060
PCRC16 bitwise comparison with reduction:	3.85280
PCRC16 bitwise comparison with task parallelism	1.51222
PCRC32	20.23205
PCRC32 with reduction	20.01593
PCRC32 with task parallelism	8.46386
PCRC32 bitwise comparison	1.78715
PCRC32 bitwise comparison with reduction:	1.86166
PCRC32 bitwise comparison with task parallelism	0.71914

B. $N = 2^{16}$. $BlockSize = 128$. $StreamDim = 4$.
 $SegSize = N/StreamDim$

Test PCRC 128 blocksize	Speedup
PCRC8	34.63720
PCRC8 with reduction	37.41059
PCRC8 with task parallelism	28.19650
PCRC8 bitwise comparison	2.24552
PCRC8 bitwise comparison with reduction:	2.70207
PCRC8 bitwise comparison with task parallelism	2.02773
PCRC16	34.14637
PCRC16 with reduction	40.48397
PCRC16 with task parallelism	25.37525
PCRC16 bitwise comparison	3.57066
PCRC16 bitwise comparison with reduction:	3.60613
PCRC16 bitwise comparison with task parallelism	1.49754
PCRC32	20.41607
PCRC32 with reduction	21.54121
PCRC32 with task parallelism	8.51250
PCRC32 bitwise comparison	1.89645
PCRC32 bitwise comparison with reduction:	1.85455
PCRC32 bitwise comparison with task parallelism	0.79350

C. $N = 2^{16}$. $BlockSize = 256$. $StreamDim = 4$.
 $SegSize = N/StreamDim$

Test PCRC 256 blocksize	Speedup
PCRC8	35.11670
PCRC8 with reduction	37.36065
PCRC8 with task parallelism	29.85230
PCRC8 bitwise comparison	2.66612
PCRC8 bitwise comparison with reduction:	2.71714
PCRC8 bitwise comparison with task parallelism	2.00425
PCRC16	33.09628
PCRC16 with reduction	39.25975
PCRC16 with task parallelism	27.58368
PCRC16 bitwise comparison	3.53824
PCRC16 bitwise comparison with reduction:	3.66416
PCRC16 bitwise comparison with task parallelism	1.54122
PCRC32	20.29566
PCRC32 with reduction	21.39834
PCRC32 with task parallelism	7.98474
PCRC32 bitwise comparison	1.75776
PCRC32 bitwise comparison with reduction:	1.80291
PCRC32 bitwise comparison with task parallelism	0.68456

D. $N = 2^{16}$. $BlockSize = 128$. $StreamDim = 2$.
 $SegSize = N/StreamDim$

Test PCRC 128 blocksize	Speedup
PCRC8	21.32339
PCRC8 with reduction	39.41265
PCRC8 with task parallelism	32.97171
PCRC8 bitwise comparison	2.70535
PCRC8 bitwise comparison with reduction:	2.64768
PCRC8 bitwise comparison with task parallelism	2.48787
PCRC16	33.66302
PCRC16 with reduction	38.77745
PCRC16 with task parallelism	30.44083
PCRC16 bitwise comparison	3.46636
PCRC16 bitwise comparison with reduction:	3.59283
PCRC16 bitwise comparison with task parallelism	2.80089
PCRC32	20.42403
PCRC32 with reduction	21.38471
PCRC32 with task parallelism	7.35099
PCRC32 bitwise comparison	1.87434
PCRC32 bitwise comparison with reduction:	1.86528
PCRC32 bitwise comparison with task parallelism	0.72902

E. $N = 2^{16}$. $BlockSize = 128$. $StreamDim = 8$.
 $SegSize = N/StreamDim$

Test PCRC 128 blocksize	Speedup
PCRC8	35.05756
PCRC8 with reduction	38.59024
PCRC8 with task parallelism	11.97966
PCRC8 bitwise comparison	2.65855
PCRC8 bitwise comparison with reduction:	2.73349
PCRC8 bitwise comparison with task parallelism	1.20074
PCRC16	34.06469
PCRC16 with reduction	38.96202
PCRC16 with task parallelism	11.68695
PCRC16 bitwise comparison	3.87572
PCRC16 bitwise comparison with reduction:	3.22093
PCRC16 bitwise comparison with task parallelism	1.58960
PCRC32	20.83102
PCRC32 with reduction	22.14467
PCRC32 with task parallelism	6.38649
PCRC32 bitwise comparison	1.94807
PCRC32 bitwise comparison with reduction:	1.91344
PCRC32 bitwise comparison with task parallelism	0.65893

F. $N = 2^{16}$. $BlockSize = 128$. $StreamDim = 4$.
 $SegSize = N/8$

Test PCRC 128 blocksize	Speedup
PCRC8	35.41841
PCRC8 with reduction	40.88182
PCRC8 with task parallelism	17.13281
PCRC8 bitwise comparison	2.40438
PCRC8 bitwise comparison with reduction:	2.97784
PCRC8 bitwise comparison with task parallelism	0.47928
PCRC16	34.66451
PCRC16 with reduction	42.29803
PCRC16 with task parallelism	16.70620
PCRC16 bitwise comparison	4.04358
PCRC16 bitwise comparison with reduction:	3.76574
PCRC16 bitwise comparison with task parallelism	1.51226
PCRC32	21.27053
PCRC32 with reduction	21.97304
PCRC32 with task parallelism	7.47568
PCRC32 bitwise comparison	1.92987
PCRC32 bitwise comparison with reduction:	1.90569
PCRC32 bitwise comparison with task parallelism	0.52763

G. NVIDIA Nsight Graphics

NVIDIA Nsight Graphics is a low overhead performance analysis tool designed to provide insights developers need to optimize their software. Unbiased activity data is visualized within the tool to help users investigate bottlenecks, avoid inferring false-positives, and pursue optimizations with higher probability of performance gains. Users will be able to identify issues, such as GPU starvation, unnecessary GPU synchronization, insufficient CPU parallelizing, and even unexpectedly expensive algorithms across the CPUs and GPUs of their target platform.

The Figure 4 shows the standard algorithm profiling where you can see on the left the kernel execution rates compared to memory transfer.

In the Figure 5 you can see how this time kernel execution prevails because there is less divergence.

With the Figure 6 and 7 you can see that execution is split over multiple streams and that in some cases, in the timeline, kernel execution and memory transfer occurs in parallel.

IX. CONCLUSIONS

This paper shows a systematic method to calculate CRC in parallel using the Galois Field property. The method can be easily expanded to all feasible r parallel-input bits for fast CRC calculation without added complexity in the developing process or practical limitation like the size of lookup tables needed in other approaches.

The implementation with tasks parallelism does not give better performance because the overhead of the technique is not positively compensated on the improvement that leads to performance. So the best result with parallelism tasks is obtained for $StreamDim = 4$ and $SegSize = N/StreamDim$.

In general the best performances are with $BlockSize = 128$, while in the case of parallelism tasks the best performances are with $StreamDim = 2$ and $SegSize = N/StreamDim$ using the reduction algorithm that reduces the divergence of the threads, which on average increases

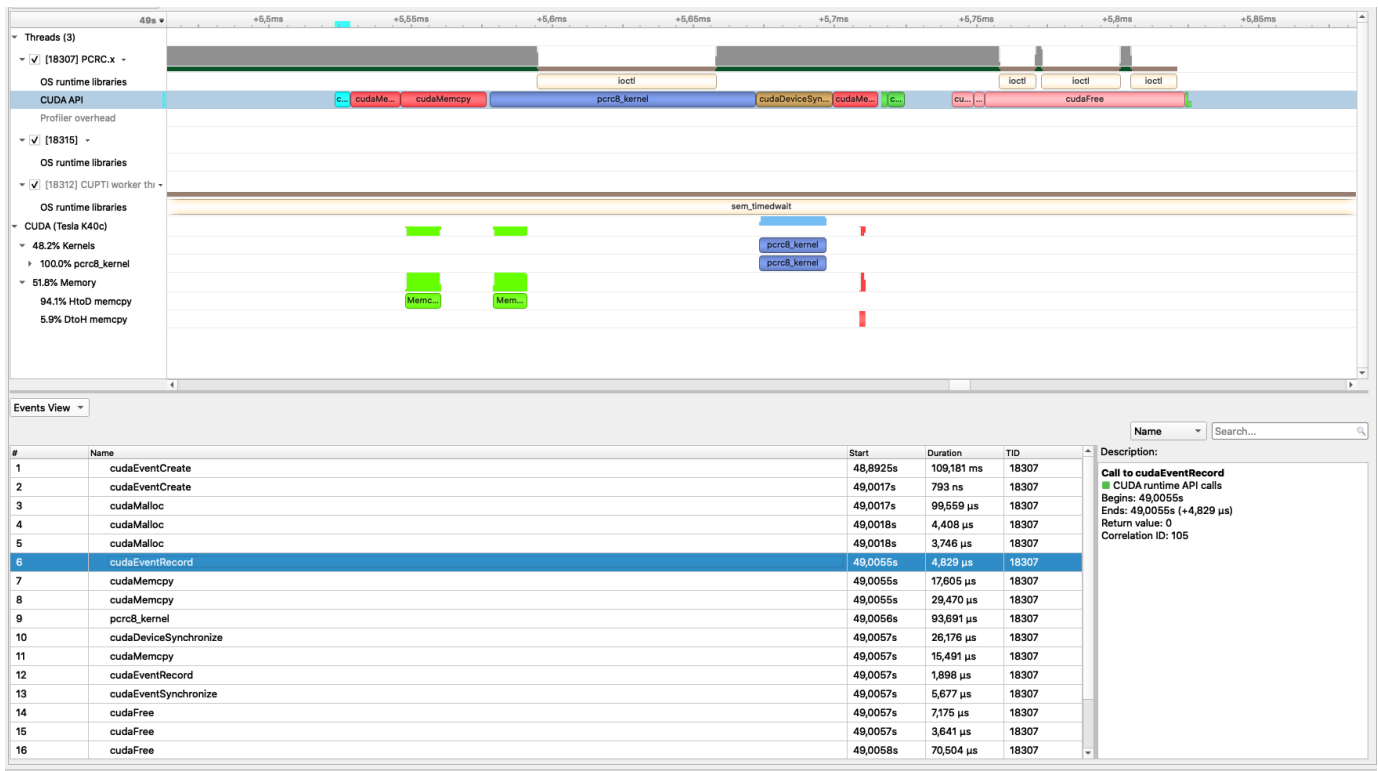


Figure 4. NVIDIA NSIGHT GRAPHICS PCRC8

the performance compared to the standard solution of some speedup units.

From these results we can say that the PCRC speedup is more +2.5.

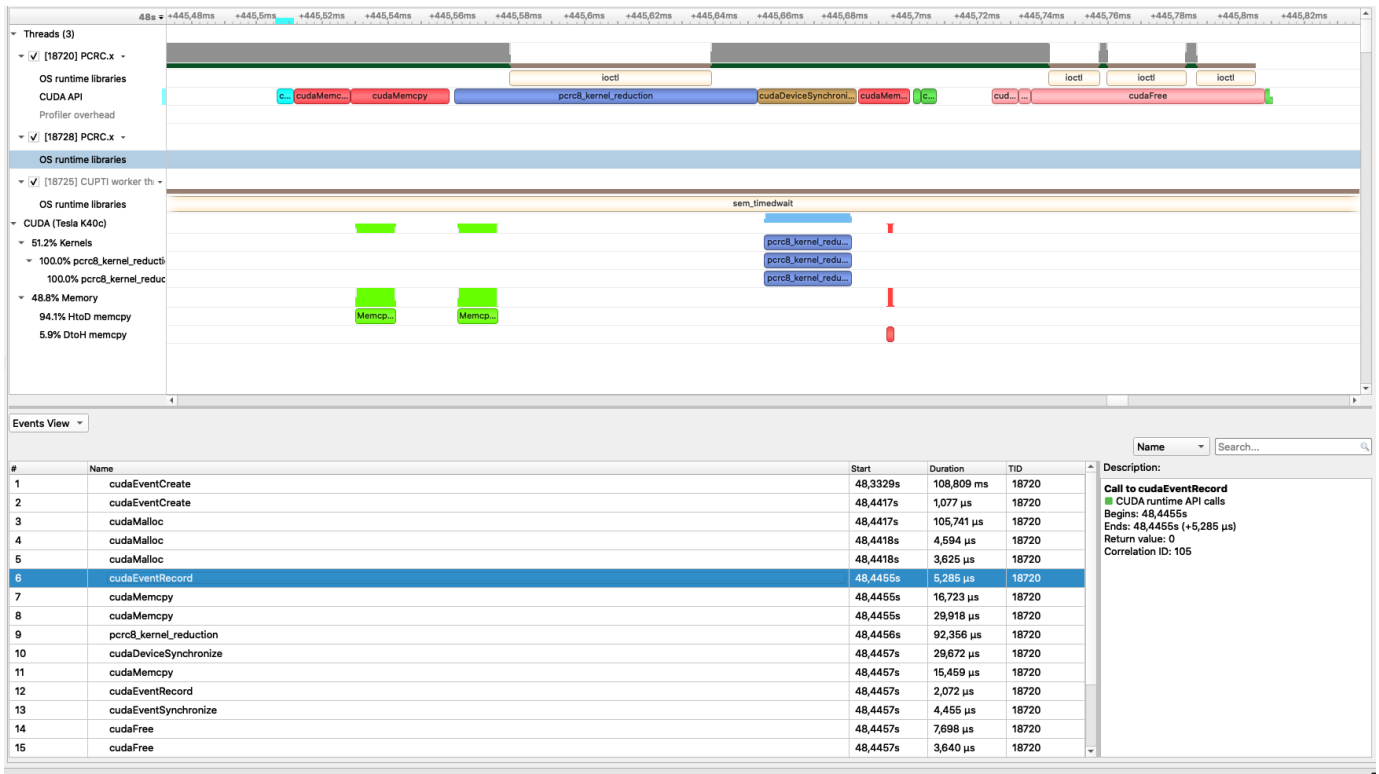


Figure 5. NVIDIA NSIGHT GRAPHICS PCRC8 REDUCTION

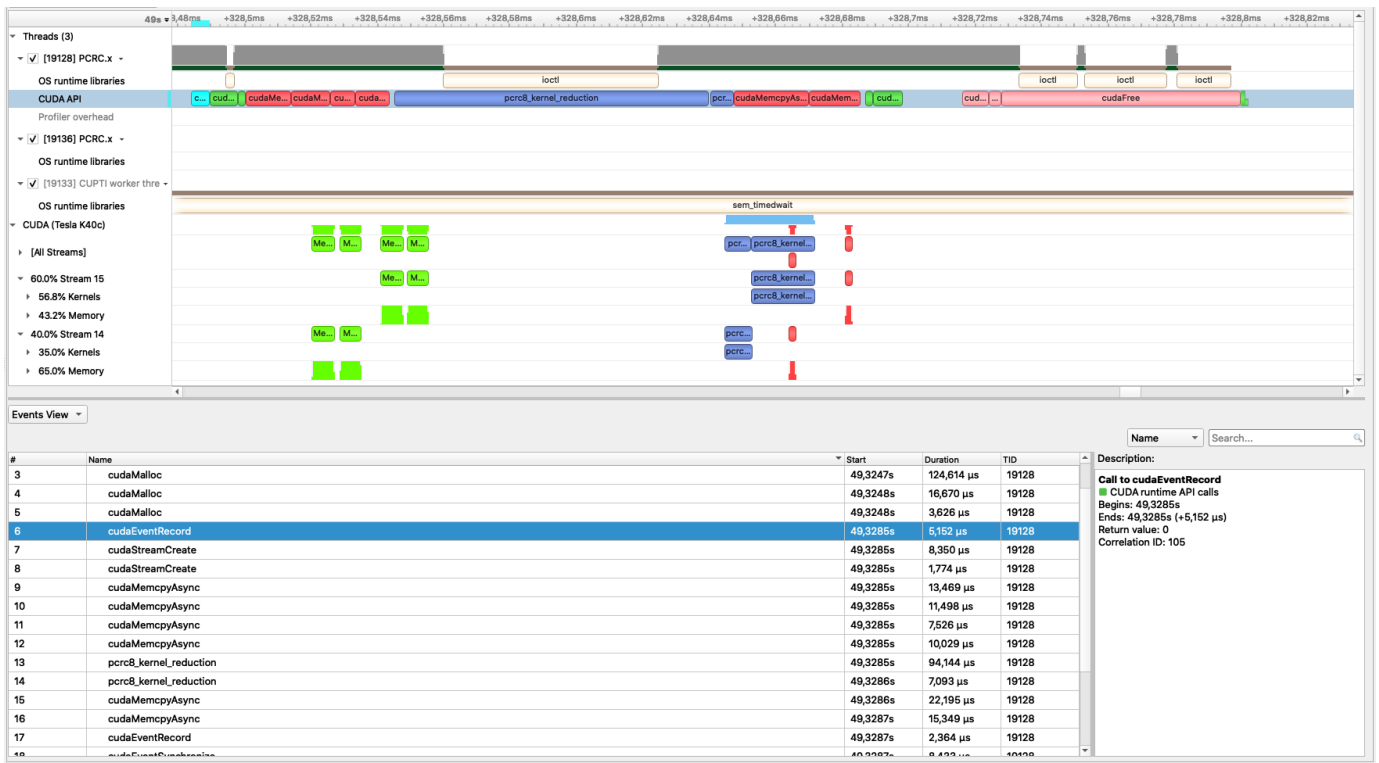


Figure 6. NVIDIA NSIGHT GRAPHICS PCRC8 TASK PARALLELISM - CASE 1




Figure 7. NVIDIA NSIGHT GRAPHICS PCRC8 TASK PARALLELISM - CASE 2